

\$4.12 of benchmark → a defensible AMD hardware business case.

\$4.12 BENCHMARK · 124 MIN · 62 PTS	256K CONTEXT · 1 MI300X (VERIFIED)	31/31 DEVS/GPU CLEAN · 8K–64K	6.49× THROUGHPUT 8K VS 32K	3.5 min DEPLOY VS 2–5 DAYS (H100)
---	--	---	--	---

AMD MI300X VS NVIDIA H100 — TCO		
LEVER	MI300X	H100
Unit price est	\$10–15K	\$32–40K+
HBM per card	192 GB	80 GB
\$ / GB VRAM (cloud)	~55% cheaper	baseline
256K single-card measured	fits (77+94 GiB)	needs 2–4× sharding
Dev seats / GPU modeled	14.5 cont · 70–140 burst	—

THE AMD FLYWHEEL (WHY IT SELLS HARDWARE)

- PoC velocity: 3.5 min vs 2–5 days → AMD runs 5–10× more PoCs/week → ~1–3 extra wins/qtr = \$5–30M/yr pipeline per account team modeled
- CUDA→ROCm agent lowers the #1 Instinct blocker (switching cost).
- AITER regression filed for AMD ROCm = ecosystem "co-developer" signal.
- Every deployment = MI300X + EPYC pull-through.

Verified on real MI300X (AMD Dev Cloud, \$1.99/GPU-hr, vLLM 0.17.1 + ROCm 7.2). AMD's own Feb-2026 blog: "serve full 256K context on a single GPU" — REPOMIND is the first open-source proof. measured = empirical · modeled = estimate on public prices.

Every company in this room — on AMD.

COMPANY · JUDGE	VECTOR	THE DEFENSIBLE \$ (MODELED ON THE MEASURED FLOOR)
AMD · Nick Ni	EARN	PoC 3.5 min → 1–3 extra wins/qtr = \$5–30M/yr pipeline/team; ~\$80–120K MI300X+EPYC per converted 8-GPU node
Oracle / OCI · Karsin, Haris	EARN	MI300X @ \$6/GPU-hr sovereign: air-gapped bank ≈ \$420K/yr/node → \$4–8M/yr @10 customers
Ameren · Bhanu	SAVE	NERC-CIP \$1.54M/day/violation → auto signed audit evidence; inference must be on-prem MI300X+EPYC
Netflix · Rajat Shah	SAVE	Memory-bound: 16×H100 → ~7×MI300X; double-digit-% off a ~\$5M/yr slice (conditional on the port)
Meta · Sriharsha	SAVE	Second-source on ~600K H100-equiv = nine-figure leverage; 8K density cuts 30K-dev deploy 311→48 GPU
Apple · Suneeth	UNLOCK	~10K restricted engineers at \$0 AI → unlock (on Apple Silicon, not AMD — honest)
AWS / Amazon · ×4	EPYC	EPYC-only (no Instinct on AWS); appliance for the air-gapped segment Amazon Q can't reach; breakeven vs Cursor 3–6 mo
Google · Ian (Gemma)	MINDSHARE	Gemma load-bearing triage, air-gapped/audited = the regulated production reference + \$2K prize

lablab.ai · Pawel, Andrea SALES Their words: "one of the standout submissions... \$30B market SaaS can't touch" → cohort + sponsor pipeline

measured floor (\$4.12 · 256K on 1 MI300X · 31/31 · 6.49×) makes every **modeled** \$ figure credible. Honest: TAM led at conservative \$5–10B (\$1–2B recoverable/enterprise, task-level speedup) — the "\$30B" is lablab's public words, quoted. No % -of-salary math.