

The one-line pitch: I turned **\$4.12 of MI300X benchmark time** into the reference workload that makes AMD's **\$/GB advantage close enterprise deals** — and shipped the CUDA→ROCm agent that lowers the one thing keeping buyers on NVIDIA.

THE MEASURED FLOOR — REAL MI300X, AMD DEV CLOUD @ \$1.99/GPU-HR

\$4.12
TOTAL BENCHMARK
124 MIN · 62 PTS

256K
CONTEXT ON
ONE MI300X

31/31
DEVS/GPU CLEAN
8K-64K

6.49x
THROUGHPUT
8K VS 32K

3.5 min
DEPLOY VS
2-5 DAYS H100

28/28
TESTS PASS
+ CI GREEN

WHY THIS SELLS INSTINCT — THE AMD BUSINESS CASE

LEVER	THE NUMBER	WHERE IT LANDS FOR AMD
PoC velocity	3.5 min vs 2-5 days → 5-10x more PoCs/week	~1-3 extra enterprise wins/qtr = \$5-30M/yr incremental pipeline per account team modeled
Hardware pull-through	Each converted 8-GPU node ≈ \$80-120K MI300X+EPYC est	Every REPOMIND deployment is an Instinct + EPYC sale, not a software license
\$/GB moat, made usable	192 GB/card · ~55% cheaper/GB · 256K single-card measured	Proves AMD's own "256K on one GPU" blog claim — first open-source proof
Switching cost ↓	CUDA→ROCm agent: 39 rules + kernel-launch rewrite, git-apply-clean diff	Directly attacks the #1 Instinct adoption blocker (port pain)
Co-developer signal	AITER regression (137/144 broken) filed for AMD ROCm	Ecosystem contributor, not just a user — a story AMD DevRel amplifies

THE WEDGE NVIDIA CAN'T TAKE
Compliance-locked buyers (utilities, banks, defense, Apple-class IP) **legally cannot** use cloud coding tools. Their inference **must** be on-prem — which means Instinct + EPYC by default.

REPOMIND is air-gapped, with an **Ed25519-signed, hash-chained audit log** — turning a regulator's requirement into AMD hardware demand.

PROOF IT'S REAL, NOT A DEMO
MIT-licensed. **make demo-offline** runs with no GPU/keys (mock-vLLM). Docker + GitHub Actions CI. 28 passing tests.

lablab.ai (official): *"a serious technical argument for a \$30B market SaaS legally cannot touch... one of the standout submissions."*

THE ASK

- **Kickoff MI300X credits** — one fresh live run so the ACT II demo/video is on current hardware (last was ACT I).
- **Intro to one AMD enterprise account team** — I'll be the reference workload in a real MI300X PoC (utility / bank / sovereign).
- **DevRel amplification** of the "VerifiedROCm" open benchmark + the AITER regression fix — positions AMD as the open, honest-benchmark platform.

measured = empirical on real MI300X (vLLM 0.17.1 + ROCm 7.2, single card, Qwen3-Coder-Next-FP8 80B/3B MoE) · **modeled** = estimate on public hardware prices. Per-org \$ figures are modeled on the measured 14.5-140 seats/GPU + public pricing; the \$4.12 / 256K / 31-dev floor is what makes them credible. Sardor Razikov · AMD ACT I winner (1st AI Agents + Outstanding Social Engagement) · razikovsardor1@gmail.com