

REPO**MIND** v3

The on-prem AI coding agent + **LLM cost-router**. For the engineers **legally barred** from cloud AI — and the enterprises **bleeding** on it. One engine. **100% on AMD**.

-89%

TOKENS · MEASURED

-90%

CACHED INPUT · AMD-SERVED

\$4.12

ACT I · 256K REPO · MI300X

\$8.4B

/YR LLM-API MARKET IT ROUTES¹⁰

Barred from cloud AI — or bleeding on it.

Path A · Locked out

REGULATED ENTERPRISES LEGALLY CAN'T USE CLOUD AI CODING

Samsung — banned ChatGPT after source-code leak ¹	May 2023
JP Morgan — restricted ChatGPT firm-wide ²	Feb 2023
Apple — blocked ChatGPT and GitHub Copilot ³	2023
US Space Force — banned generative AI / CUI ⁴	2023
Healthcare — PHI can't leave; avg HIPAA fine ⁵	~\$1.2M
EU AI Act — non-compliance penalty ⁶	≤€35M / 7%

\$0 AI

50,000+ regulated engineers with nothing. On-prem is the only class of tool they can deploy.

Path B · Token tax

EVERYONE ELSE IS DROWNING IN RUNAWAY AI COST

Uber — burned its 2026 AI-coding budget in 4 months, capped \$1,500/mo ⁷	4 mo
GitHub Copilot — usage billing: \$29 → \$750/mo per dev ⁸	~25x
Cursor — \$2B ARR, negative margin, built own router ⁹	\$2B ARR
Enterprises — spending >\$250K/yr on LLM APIs ¹⁰	37%
Model-API spend — doubled YoY, and climbing ¹⁰	\$8.4B/yr
Microsoft — cut internal Claude Code licenses over cost ¹⁸ (Walmart·Amazon·Cisco capping too)	Jun 2026

10–50x

"The subsidy era ended." Cost is now everyone's problem — the banned worst of all.

One engine covers both: on-prem for the barred, a cost-router for the bleeding — and both flow to AMD.

On-prem agent + cost-router. Every call is **AMD compute.**

WHAT IT IS

A **SC-TIR agent** (Plan→Call→Observe→Think→Answer) reasoning over your codebase **on-prem**, **air-gapped**. The same engine **routes/proxies** existing Claude/OpenAI/Gemini spend and cuts it.

MEASURED COST LAYER · LIVE ON MACBOOK + AMD FIREWORKS

−89% AST skeleton · **−90%** cached input¹³ · **cascade skips** the deep model when the cheap tier is sure. Industry routing confirms **40–85%** cuts, ~74% of queries never need the frontier¹². App streams **\$ saved per query**.

TRUST — THE AI FLIGHT RECORDER

Ed25519-signed, hash-chained, replayable audit of every action — auditor-legible (EU AI Act Art. 12⁶). Compliance becomes automatic evidence.

100% AMD — NO CLOSED MODELS

Runs on **Fireworks (AMD Instinct-powered)** + **Radeon → MI300X** on-prem. **No OpenAI / Anthropic / Gemini**. Every deployment is an **AMD pull-through** in the segment NVIDIA's CUDA moat can't defend.

GEMMA IS THE OPEN-WEIGHT TIER — IN GOOGLE'S OWN WORDS¹⁶

"Purpose-built for **agentic workflows**... power your **IDEs, coding assistants**... highest **compliance guarantees for regulated workloads**... integrate with **AMD GPUs via ROCm**."

That is REPOMIND, shipped. Gemma markets regulated-on-prem-on-AMD but has **no named production reference** — we're it.

IT ACTUALLY RUNS

docker compose up · 38 passing tests + CI · **make demo-offline** (no GPU/keys) · CUDA→ROCm · path-sandboxed.

Millions saved per deployment.

An **\$8.4B** flow, funneled to AMD.

A 5,000-dev bank recovers **\$11–22M/yr** in locked coding productivity — even at 1% capture that's \$1.5M, 8× the \$180K license, payback in days. On-prem open-weight = **\$0 per-token** vs banned Copilot/Cursor at \$1.1–2.3M/yr forever.

PATH A — ON-PREM SAVINGS (BANNED ENTERPRISES)

ENTERPRISE	DEVS	AMD GPUS	SAVE/YR ^A
JP Morgan	50,000	~77 MI300X	~\$7.6M
Meta	30,000	~48	~\$4.6M
Apple	10,000	~16	~\$1.5M
Lockheed	20,000	~32	~\$3.1M

PATH B — ROUTER SAVINGS (BLEEDING ENTERPRISES)

PAYER	TODAY	AFTER –40–70% ^B
Uber (~5,000 eng)	~\$30M/yr	save \$12–21M
Copilot dev	\$750/mo	save \$300–525
\$1M/yr spender	\$1M/yr	save \$400–850K

PRICING · HIGH-MARGIN OPEN-CORE

\$500 · \$2,500 · \$15K / mo (starter → air-gapped) + AMD-cloud burst. MIT open-core → **\$180K–2M/yr** enterprise. No per-token COGS on-prem.

SERVICEABLE MARKET — BOTH PATHS

Path A: 2,000+ regulated × \$180K–2M = **\$0.4–4B**. **Path B:** >\$250K/yr cohort ≈ \$3B flow → ~\$1.5B savings pool → **\$225–300M** capture ceiling.

THE NESTED MARKET STACK — SOURCED

AMD AI-accelerator TAM by 2030 ¹⁴	\$1 TRILLION
Enterprise-LLM market by 2034 ¹¹	\$71B
Model-API spend now (measured) ¹⁰	\$8.4B/yr
Routing cut, no quality loss ¹²	40–85%

^a Modeled on the measured 6.49× throughput + ~31 devs/GPU — AMD adoption from a **\$0 baseline, not fewer sales**. ^b Modeled on measured levers + industry routing¹²; company \$ ranged, not booked. Productivity recovery (hammer) is separate and larger.

One deployment. The whole ecosystem profits.

Proof, already shipped: AMD ACT I — 1st Place AI Agents + Outstanding Social Engagement (two trophies · 481 submissions · 10,770 builders) · verified on real MI300X **\$4.12 / 256K / single card** · AMD Developer officially featured REPOMIND · 3 Zenodo preprints. **lablab.ai:** *"one of the standout submissions — the business case is airtight."*

AMD — THE PLATFORM

Q1 FY2026 Data Center \$5.8B, +57% YoY¹⁵;

Lisa Su names **agentic AI** as the driver — REPOMIND is one. Every deployment = **\$200–500K** Instinct+EPYC pull-through. Agents drive **CPU** demand too — EPYC just hit a record **46.2% server-CPU share**¹⁷. Feeds the **\$1T-by-2030** accelerator TAM¹⁴.

FIREWORKS AI

The paid burst tier **is Fireworks** (on AMD Instinct). Cheaper burst → more enterprises adopt → **more total AMD-hosted consumption**, per customer, recurring. The router sends the cheap 60–80% here and on-prem.

GEMMA / DEEPMIND

A live **regulated on-prem Gemma-on-AMD** reference Google's marketing lacks¹⁶ — the champion demo Ian Ballantyne already evangelizes ("Gemma offline for local coding"). Targets the **Best-Use-of-Gemma** bonus.

ENTERPRISE

\$11–22M/yr recovered (Path A) or **40–70% bill cut** (Path B) · \$0 per-token on-prem · IP-leak + EU-AI-Act risk down. Payback in **days**.

Sources — **1** Forbes '23 (Samsung) · **2** CNN '23 (JPM) · **3** TechCrunch '23 (Apple/Copilot) · **4** Air&Space Forces '23 (Space Force/CUI) · **5** Brellium/HHS-OCR (HIPAA ~\$1.2M) · **6** artificialintelligenceact.eu (Art. 99, ≤€35M/7%) · **7** TechCrunch '26 (Uber \$1,500 cap) · **8** GitHub Blog / TechCrunch '26 (Copilot \$29→\$750) · **9** TechCrunch/Reuters '26 (Cursor \$2B ARR, thin margin) · **10** Menlo Ventures '25 (\$3.5B→\$8.4B API spend; 37%>\$250K) · **11** Global Market Insights (\$8.8B→\$71.1B, 26% CAGR) · **12** Orq/TrueFoundry/IBM '26 (routing 40–85%) · **13** Anthropic/OpenAI/Google docs (cache –90%) · **14** CNBC — Lisa Su, AMD Financial Analyst Day Nov '25 (\$1T by 2030) · **15** CNBC/DataCenterDynamics (AMD Q1 FY26 DC \$5.8B, +57%) · **16** blog.google — Gemma 4 (agentic / coding / regulated / AMD ROCm) · **17** AMD "Agentic AI Changes the CPU/GPU Equation" / Mercury Research — EPYC record 46.2% server-CPU revenue share; server-CPU TAM →\$120–170B by 2030 · **18** thenextweb / cryptobriefing '26 — Microsoft pulls internal Claude Code licenses over cost (\$500–2,000/eng/mo).

Creativity · Product/Market · Completeness · Use of AMD — all four, 100% AMD, measured & honest. Company \$ = modeled on cited public data, ranged, not booked; named companies are not customers.