

50,000 engineers. \$0 AI today. One legal path — on AMD.

Regulated enterprises are **legally barred** from cloud AI coding — their engineers have **nothing**. On-prem, open-weight on AMD is the only path. **The conservative hammer:** one 5,000-dev bank recovers **\$11–22M/yr** in coding productivity — **even at 1% capture that's \$1.5M, still 8× the \$180K license, payback in days**. The math needn't be exact — it's directionally undeniable.

\$4.12 256K REPO · 1 MI300X MEASURED	31 DEVS/GPU CLEAN · 8K–64K MEASURED	6.49× THROUGHPUT 8K VS 32K MEASURED	–89/–90% SKELETON / CACHED INPUT MEASURED	\$0 PER-TOKEN ON-PREM (OPEN-WEIGHT)
---	--	--	--	---

REGULATED ENTERPRISE	DEVELOPERS	ON-PREM AMD GPUS¹	GPU-EFFICIENCY SAVING/YR²	+ ON-PREM INFERENCE
JP Morgan · banking	50,000	~77 MI300X (vs 500 naïve)	~\$7.6M/yr	Open-weight runs on the bank's own AMD GPU → \$0 per-token , data never leaves. The paid Fireworks burst is cut –89/–90% by the cost layer. vs SaaS they can't legally use: Copilot/Cursor ≈ \$19–39/seat/mo = \$1.1–2.3M/yr forever, and banned.
Meta · 30K devs	30,000	~48 MI300X (vs 311)	~\$4.6M/yr	
Apple · iOS eng	10,000	~16 MI300X (vs 104)	~\$1.5M/yr	
Lockheed · cleared	20,000	~32 MI300X (vs 210)	~\$3.1M/yr	
Netflix · senior ML	3,000	~7 MI300X (vs 42)	~\$0.6M/yr	
Illustrative 5-account slice	113,000	~180 MI300X	~\$17.4M/yr modeled	+ ~\$0 marginal inference on-prem

Efficiency (**measured** 6.49×) makes on-prem AMD viable for a segment that otherwise deploys **zero** — so this is AMD adoption **from a \$0 baseline, not fewer sales**. GPU counts modeled on the 6.49× + ~31 devs/GPU at realistic peak concurrency; efficiency saving vs a naïve 32K deployment @ ~\$1.99/GPU-hr. Productivity recovery (task-level speedup) is **separate and larger** (see hammer above).

measured

 = ACT I on real MI300X + live bench ·

modeled

 = public prices. Illustrative of the compliance-locked segment, not customers.

One deployment. The whole ecosystem profits.

AMD — RAMINE ROSEN · MALINI · NAG LOHITH

EARN. Every on-prem deployment = **\$200K–\$500K** Instinct + EPYC pull-through modeled. 3.5-min PoC (measured) → **5–10× more PoCs/week** → 1–3 extra wins/qtr = **\$5–30M/yr pipeline** per account team. REPOMIND proves the cost-down economics that make on-prem viable — AMD sales from a **\$0 baseline**, a segment NVIDIA is **legally locked out of**.

FIREWORKS AI — SPONSOR · AMD-HOSTED

EARN. The paid burst tier **is Fireworks**. Even minimized (–89/–90%), every enterprise customer streams tokens through Fireworks for the hard queries → **recurring AMD-hosted API revenue**. Cheap-enough burst = more enterprises adopt = **more total consumption**, not less.

GOOGLE DEEPMIND / GEMMA — IAN BALLANTYNE

MINDSHARE. Gemma is the open-weight router tier that runs on the customer's AMD GPU → a **regulated on-prem production reference** Gemma marketing lacks. A live Gemma-on-AMD run unlocks the **\$2,000 "Best AMD-Hosted Gemma"** bonus.

LABLAB.AI / NATIVELYAI — PAWEL · ANDREA · ADAM

SALES. An honest, **measured** Track-3 build = a case study that drives **cohort registration + sponsor pipeline**, and validates the AMD × Natively × Fireworks stack they sell. ACT I winner → v3 = the ecosystem continuity story.

The chain: a regulated enterprise recovers **\$11–22M/yr** locked productivity + kills \$100M–\$1B IP-leak + ≤€35M EU-AI-Act risk → pays REPOMIND **\$180K/yr** → which **requires AMD** (\$200–500K) + **feeds Fireworks** (burst) + **references Gemma** + gives **lablab/Natively** the case study. **Every link is AMD-served, open-weight, cost-down** — exactly what Malini, Ramine and Nag are looking for.

REPOMIND the business: MIT open-core → enterprise license **\$180K–\$2M/yr**, high-margin (open-weight — no model-training cost, no per-token COGS on-prem). Serviceable segment ≈ **2,000+ regulated enterprises × \$180K–2M = \$0.4–4B** modeled bottoms-up (banks · defense · healthcare · IP-sensitive big-tech).

All per-org \$ figures are modeled on the measured floor + public prices, ranged, never asserted as booked. Productivity = task-level speedup on the coding share of dev time. Gemma bonus conditional on a live Gemma-on-AMD run.