


We built the cost engine for the banned.

The market **handed us everyone.**



\$4.12 256K REPO · MI300X MEASURED	–89/–90% SKELETON / CACHED INPUT MEASURED	40–85% ROUTING CUT (INDUSTRY SOURCED)	\$8.4B→\$71B LLM-API MARKET '25→'34	→ AMD EVERY OPTIMIZED CUSTOMER
---	--	---	---	--

15–60× between tiers. 74% never needed the frontier.

PRICE SURFACE · PER 1M TOKENS · JUL 2026

MODEL	INPUT	OUTPUT	CACHE
Claude Opus 4.6	\$5.00	\$25.00	−90%
GPT-5.4	\$2.50	\$15.00	−90%
GPT-5.4 Pro	\$30.00	\$180.00	−90%
Gemini 3.1 Pro	\$2.00	\$12.00	−90%
Claude Sonnet 5	\$2–3	\$10–15	−90%
Gemini 3 Flash	\$0.50	\$3.00	−90%
Open-weight · Fireworks (on AMD)	~\$0.1–0.5	~\$0.2–0.8	−80/90%
Open-weight · on-prem AMD	≈ \$0	≈ \$0	—

Every query sent to a frontier model that a cheap tier could answer is pure waste. Prices fell ~80% in '25→'26 — the arbitrage only widens.

ROUTING WORKS — THIRD-PARTY VALIDATION

40–85% bill cut, no visible quality loss. A router keeps 95% of GPT-5.2 quality using it for only 26% of requests → 74% goes to cheap tiers. RouteLLM/FrugalGPT: 2×–98% at equal quality. → our modeled 40–70% is conservative.

THE PAIN, BY NAME

Uber	whole yr budget · 4 mo
Copilot org	\$29 → \$750/mo (25×)
Cursor	neg. margin → built own router

RepoMind's edge (not "we route" — everyone does): coding-native AST compression + Ed25519 Flight Recorder audit + the AMD open-weight migration path. Cut the bill today, move the cheap majority to AMD.

Case B market: \$8.4B → \$71B LLM-API spend; the >\$250K/yr cohort ≈ \$3B+/yr optimizable → ~\$1.5B savings pool → capture 15–25% = \$225–300M ceiling · 1% ≈ \$2–3M · 5% ≈ \$11–15M modeled. And every optimized customer becomes an AMD-inference customer — the funnel AMD doesn't have.

One engine covers both markets — and every link is AMD.

The AMD tailwind: Lisa Su (Financial Analyst Day, Nov 2025) sizes the AI accelerator TAM at **\$1 TRILLION by 2030** (up from \$500B/2028). AMD Q1 FY2026 Data Center = **\$5.8B, +57% YoY**; Su names **agentic AI** as the demand driver. **REPOMIND is an agentic-AI workload that converts that TAM into Instinct + EPYC pull-through** — in the one segment NVIDIA's CUDA moat can't defend.

AMD — RAMINE ROANE (CVP AI) · THE PLATFORM

EARN. His public mission = **"make AI ubiquitous across all AMD platforms."** REPOMIND puts AMD AI where NVIDIA-cloud can't go — **air-gapped enterprises on customer-owned Instinct/EPYC**, built on AMD's open stack (**ROCm 7 + vLLM**). Every deployment = **\$200–500K** silicon [modeled]; every optimized cloud customer = an AMD-inference migration. Owned-hardware \$0/token is where AMD's TCO argument is **strongest** (sidesteps thin AMD-rental availability).

GEMMA — IAN BALLANTYNE (DEEPMIND DEVREL)

REFERENCE. Gemma 4's own words: *"purpose-built for agentic workflows... power your IDEs, coding assistants... highest compliance guarantees for regulated workloads... integrate with AMD GPUs via ROCm."* That's REPOMIND verbatim. Gemma markets sovereignty but lists **no named regulated on-prem production reference** — REPOMIND is it. Ian personally demos **Gemma offline for local coding** ("Sovereign Escape Velocity"). A live Gemma-on-AMD run targets the **Best-Use-of-Gemma bonus**.

LABLAB · NATIVELYAI — PAWEL · ANDREA

ECOSYSTEM. lablab = #1 AI-builder ecosystem (**240K+ devs**). REPOMIND validates the exact stack they sell — **AMD compute x Fireworks API x open-weight Gemma** — and gives them the **ACT I winner → v3 continuity** case study that drives cohort + sponsor pipeline. Honest, measured, cited = the story they amplify.

THE ENTERPRISE — BOTH PATHS

PATH A (banned): Samsung (code-leak ban '23), JPMorgan, Apple (banned Copilot), Space Force, HIPAA health — recover locked productivity, \$0 egress, EU-AI-Act €35M/7% risk down. **PATH B (bleeding):** Uber (budget gone in 4mo, \$1,500 cap), Copilot (\$29→\$750). Cut the bill **40–70%**, migrate to AMD. All company \$ = modeled on cited public facts, ranged.

The chain: enterprise recovers/saves millions → pays REPOMIND → which **requires AMD** (Instinct+EPYC) + **runs Fireworks** (on AMD Instinct) + **references Gemma** (regulated on-prem) + gives **lablab** the case study. **Both markets, one engine, all flowing to AMD.** Sources: AMD IR/CNBC · Google Gemma-4 blog · Menlo Ventures · Forbes/CNN/TechCrunch. Judge roles per public posts; Gemma-bonus amount to be confirmed on the event page.